

Photometric Redshifts from Observed Images

...

with Boosted Decision Trees, Random Forests and CNNs

12.03.2021

David Rufer, Sebastián Gil and Qiang Wang (王强)

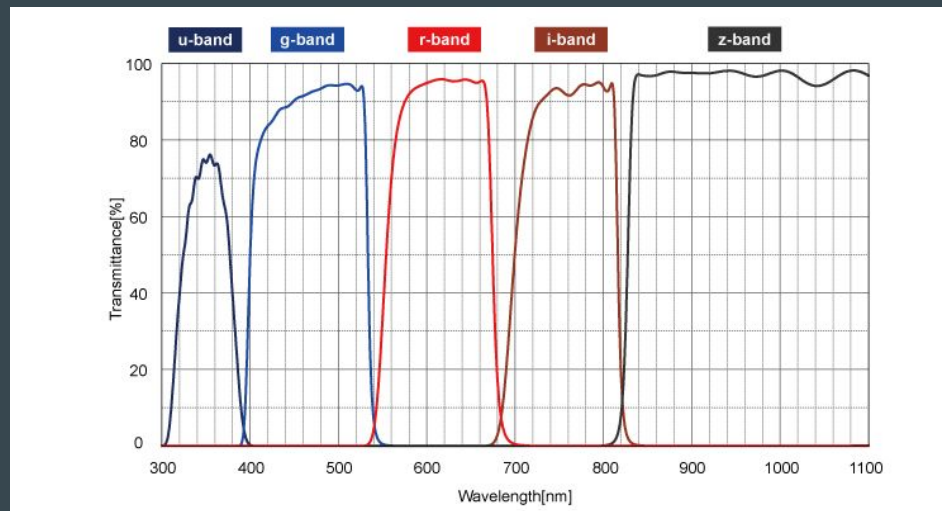
Outline

- Introduction
- Data Preparation
- Machine Learning Algorithms
- Conclusion

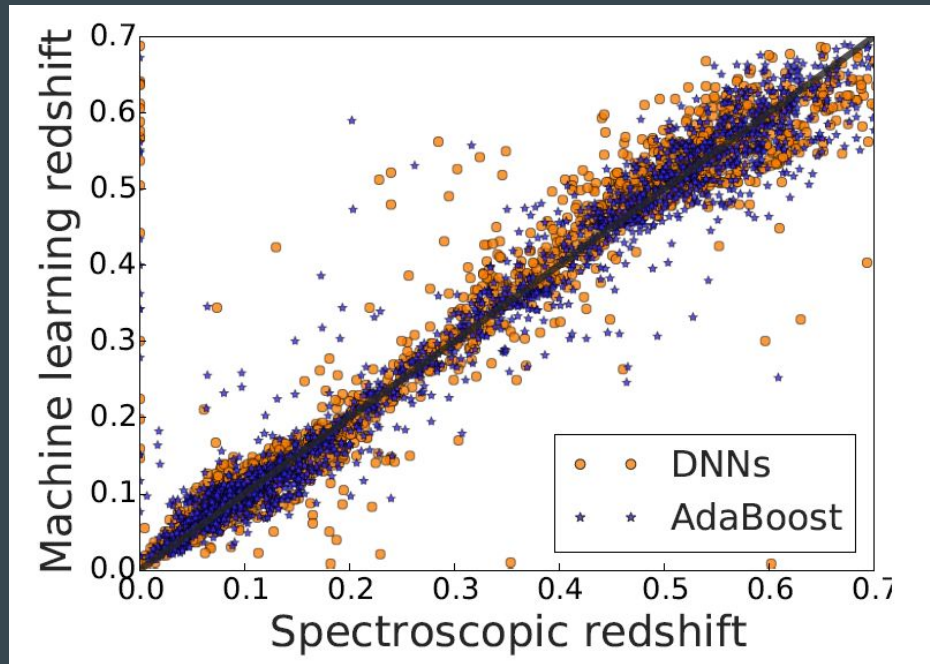
Introduction

Statement of the problem

- Redshifts
- Photometric vs Spectroscopic
- Why use machine learning?
- Why use the whole galaxy image?



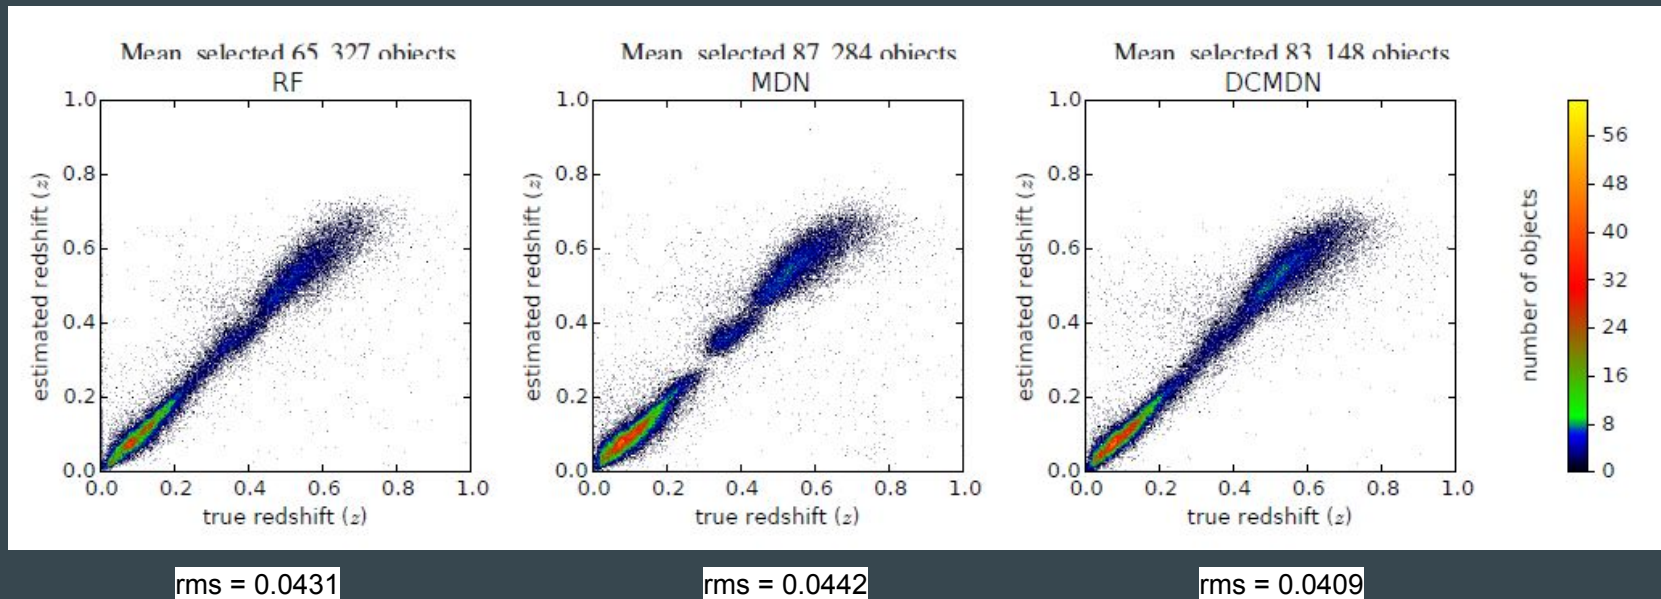
Prior Research



*“the choice of which photometric input features to train the machine architecture, from the full list of possible photometric features, is still left to the discretion of the user.... We **no longer impose our prior beliefs** upon which derived photometric features produce the best redshift predictive power.... **we completely remove the user** from the photometric redshift estimation process”.*

Hoyle, 2015


Prior Research




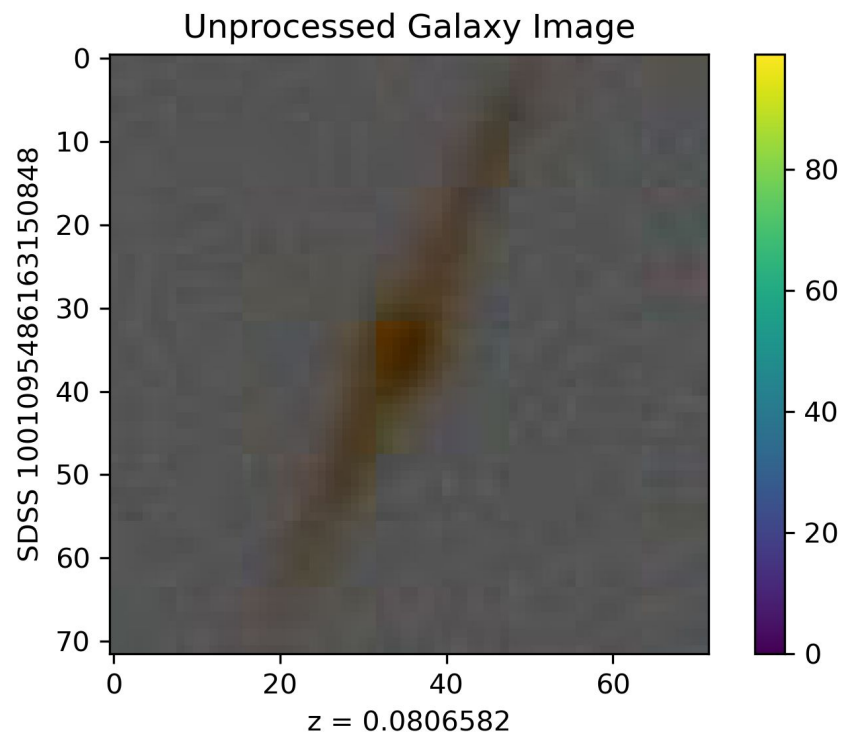
D'isanto and Polsterer, 2019




SDSS Dataset

Feature	Value
Number of galaxies	57720
Minimum redshift	1.47×10^{-7}
Maximum redshift	1.98
Average redshift	0.31
Median redshift	0.28
Mode of redshifts	0.13










Data Preparation

Structure

- Import Data and Redshifts
 - Flipping and Rotating
- Rescaling
 - Inverting
- Cropping
 - Random vs. Maximum Crop
- (Flattening)
- Splitting



Algorithms Used

Boosted Decision Trees

What are boosted decision trees?

- Combination of multiple weak learners into a single strong learner

$$F_{m+1}(x) = F_m(x) + h_m(x)$$

$$h_m(x) = -\frac{\partial L_{\text{MSE}}}{\partial F}$$

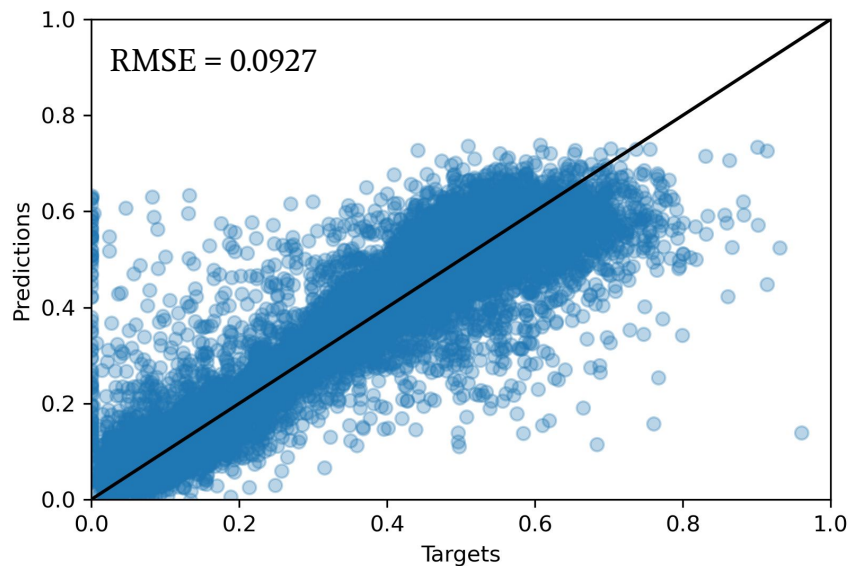
Hyperparameter Search with HyperOpt

```
space = {  
    'loss': hp.choice('loss', ['ls', 'huber', 'lad']),  
    'criterion': hp.choice('criterion', ['friedman_mse', 'mse']),  
    'learning_rate': hp.loguniform('learning_rate', np.log(0.01), np.log(0.4)),  
    'n_estimators': hp.randint('n_estimators', 150),  
    'min_samples_split': hp.randint('min_samples_split', 50),  
    'min_samples_leaf': hp.randint('min_samples_leaf', 51),  
    'max_depth': hp.randint('max_depth', 16)  
}
```

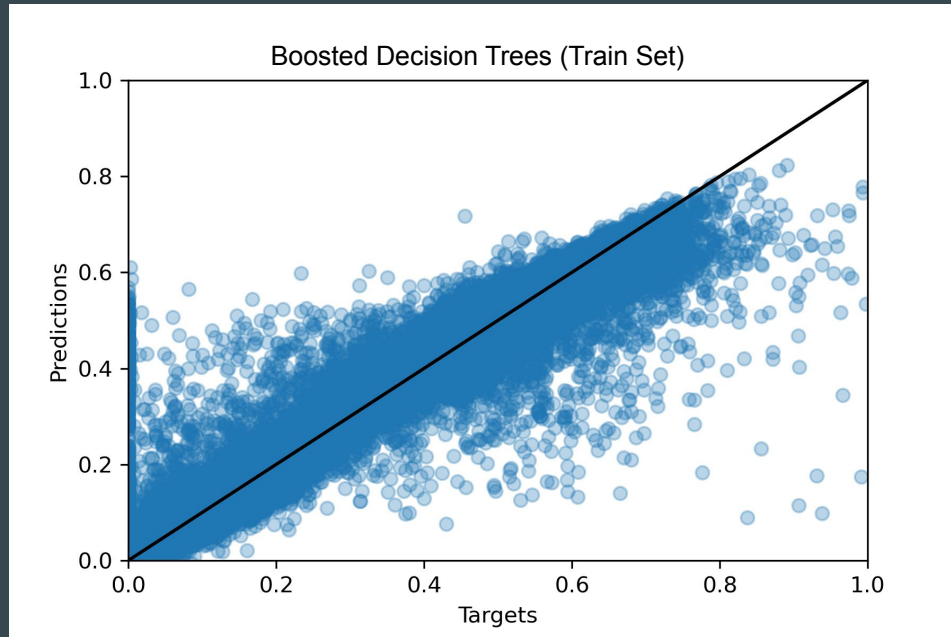
```
GradientBoostingRegressor(learning_rate=0.52, loss='huber', max_depth=14,  
                           min_samples_leaf=12, min_samples_split=43,  
                           n_estimators=77)
```

Results (change titles)

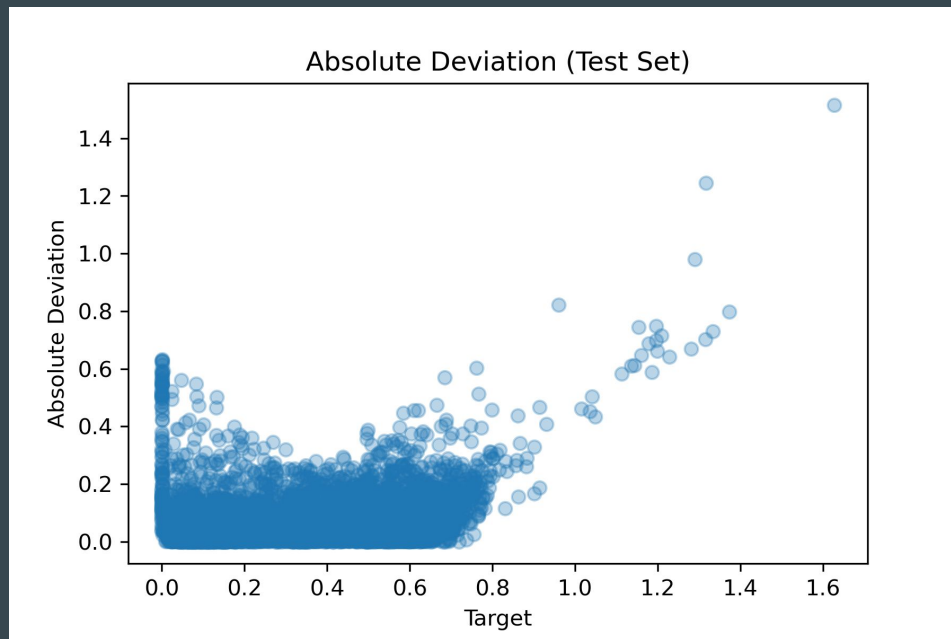
```
GradientBoostingRegressor(learning_rate=0.52, loss='huber', max_depth=14,  
                           min_samples_leaf=180, min_samples_split=650,  
                           n_estimators=77)
```



Results



Results



Convolutional Neural Network

Review of CNNs


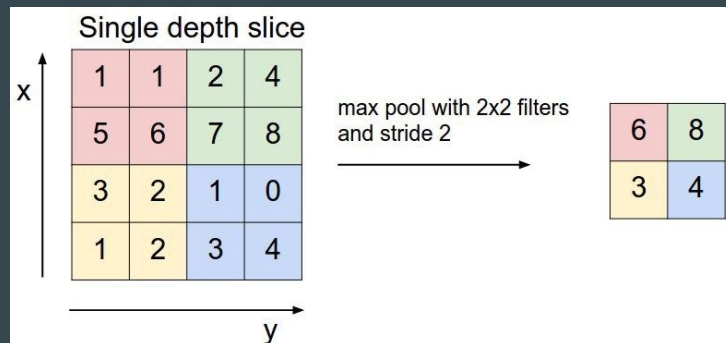



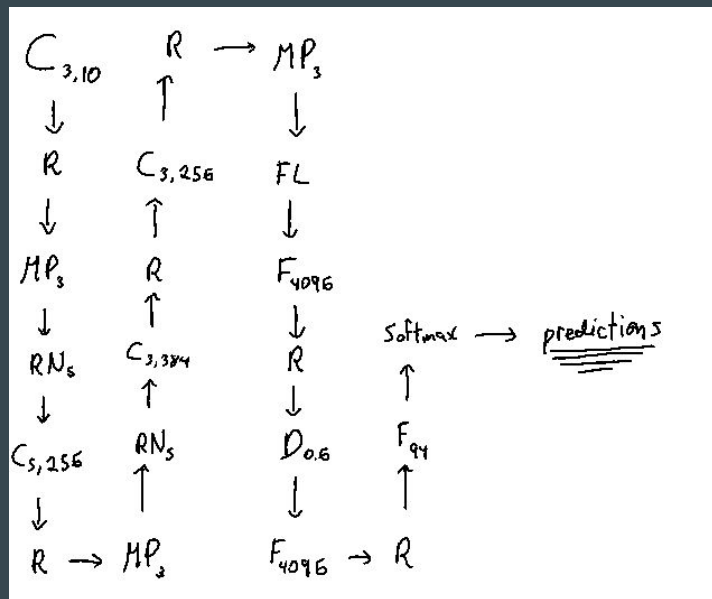
Image credit:
towardsdatascience.com



Classification vs. Regression Tasks

	Classification	Regression
targets	class list of features	array of values
loss	cross-entropy	mse, mae, logcosh
input/output	one hot output layer and labels	single output
final activation	softmax	linear

Architectures Considered




- Implemented in GraphLab (deprecated)
- Little reference \rightarrow hard to implement
- Poor performance in Keras
- Replicability...

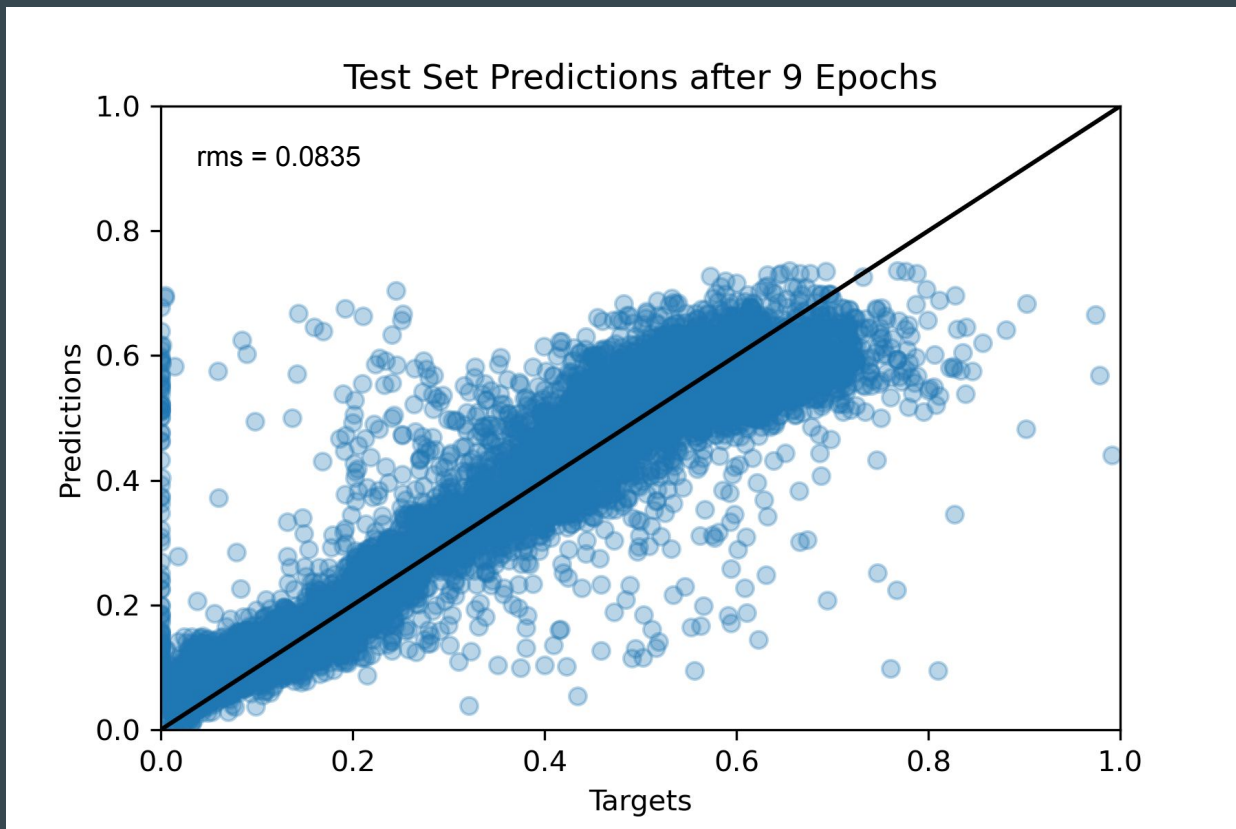
Architectures Considered


```
def createCNNModel():  
    model = Sequential()  
    model.add(Conv2D(10, (5, 5), padding='same', input_shape=(60,60,3)))  
    model.add(BatchNormalization())  
    model.add(Conv2D(32, (3, 3), activation='relu'))  
    model.add(MaxPooling2D(pool_size=(3, 3)))  
    model.add(BatchNormalization())  
    model.add(Conv2D(64, (3, 3), activation='relu'))  
    model.add(Conv2D(128, (3, 3), activation='relu'))  
    model.add(MaxPooling2D(pool_size=(3, 3)))  
    model.add(Flatten())  
    model.add(Dense(2000, activation='relu'))  
    model.add(Dropout(0.3))  
    model.add(Dense(500, activation='relu'))  
    model.add(Dropout(0.2))  
    model.add(Dense(300, activation='relu'))  
    model.add(Dense(30, activation='relu'))  
    model.add(Dense(20, activation='relu'))  
    model.add(Dense(20, activation='relu'))  
    model.add(Dense(1, activation=tf.keras.activations.linear,  
        bias_initializer=tf.keras.initializers.Constant(0.5)))  
    return model
```

Hyperparameters	Values explored
learning rate	0.1 – 0.00001
batch size	{32, 64}
loss	{'mae', 'mse', logcosh}
epochs	{4 – 16}
dropout	{0.2, 0.4, 0.6}

Results







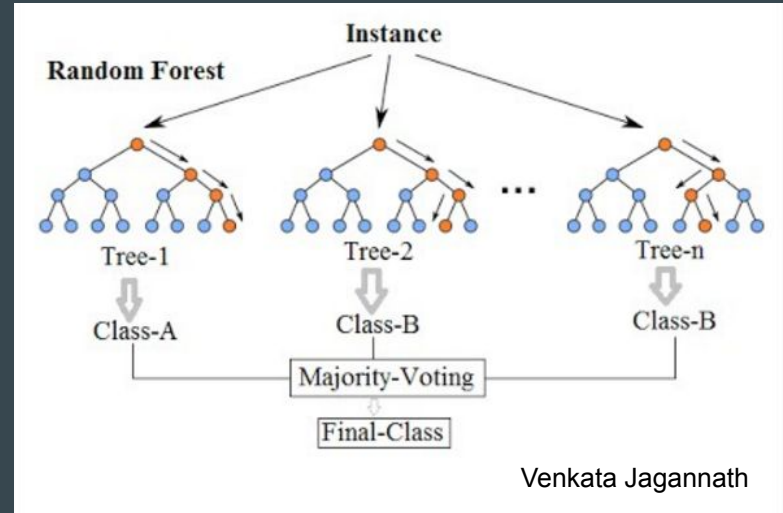
Random Forests

Random Forest Basics

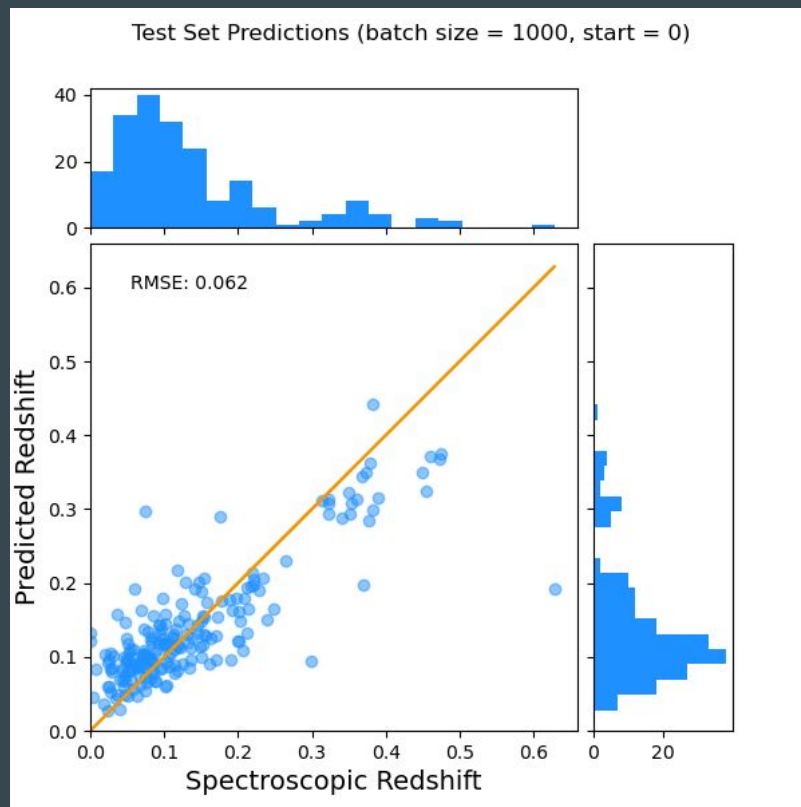
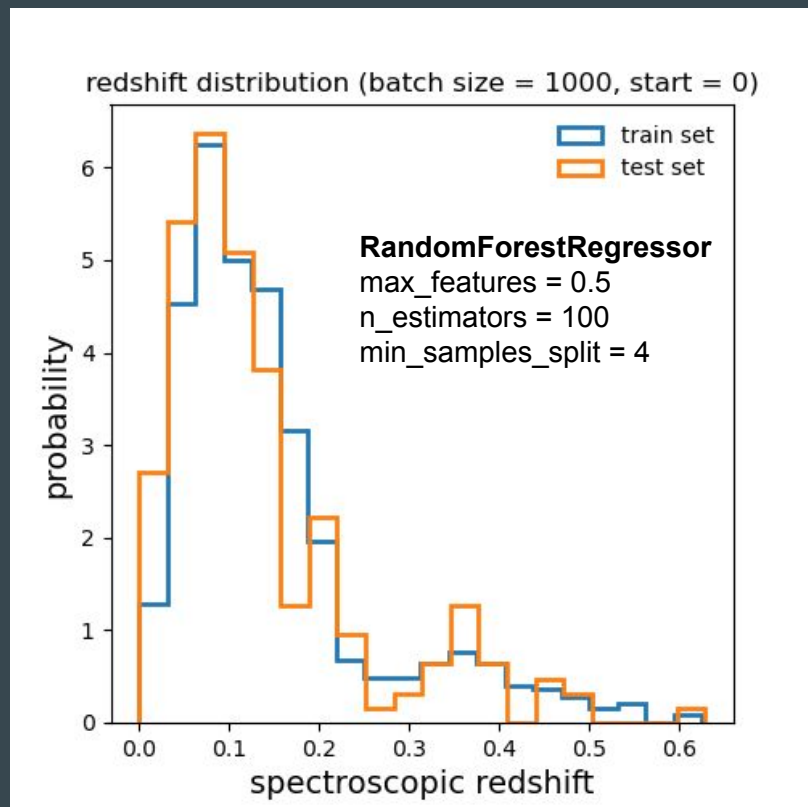
RF, aka random decision forest, is based on decision tree method and improves it by constructing a multitude of decision trees at training time.

The output is the average of the individual tree, for regression.

This method overcomes the overfitting problem with decision tree.



Random Forest: Quick Example



RandomForestRegressor: Hyperparameter tuning

- **n_estimators**: number of trees, the larger the better
- **max_features**: size of the random subsets of features to consider when splitting a node (our data has 10800 features), options: none, sqrt, float, int etc.
-
- **bootstrap** (True/False): if bootstrap samples are used when building trees. If False, the whole dataset is used to build each tree.
- **max_depth**: The maximum depth of the tree
- **min_samples_split**: the minimum number of samples required to split an internal node
- **oob_score** (True or False): whether to use out-of-bag samples to estimate the R^2 on unseen data.
- ...

Random Forest: Hyperparameter tuning with GridSearchCV

```
n_estimators = np.arange(100, 300, 2)
```

```
max_features = np.concatenate(([None, "sqrt", "log2", ], np.linspace(0.1, 0.99, 5)))
```

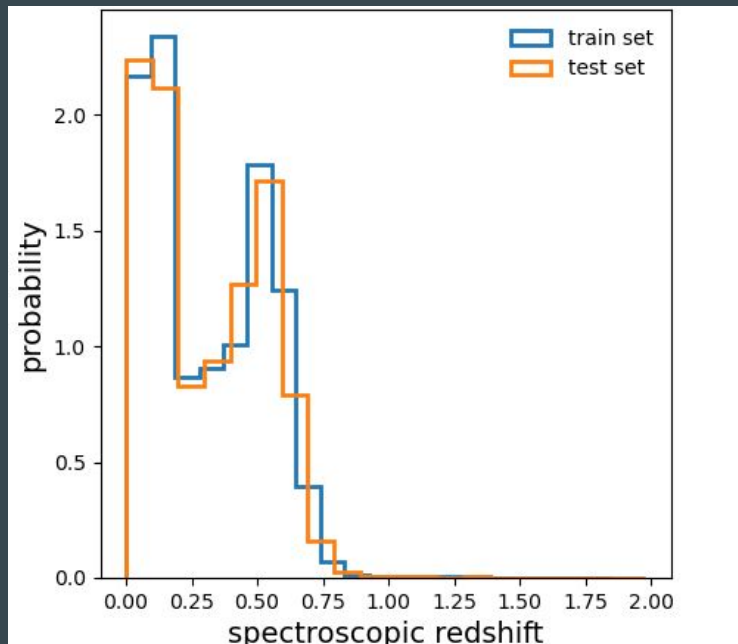
```
bootstrap = [True, False]
```

```
oob_score = [True, False]
```

```
min_samples_split = [2, 4]
```

```
max_depth = [None]
```

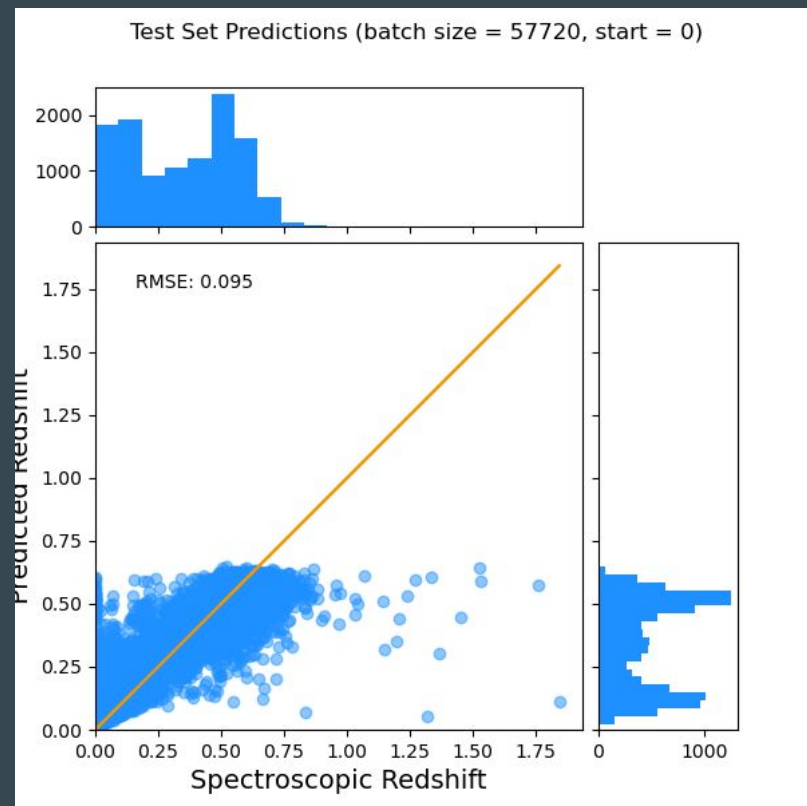
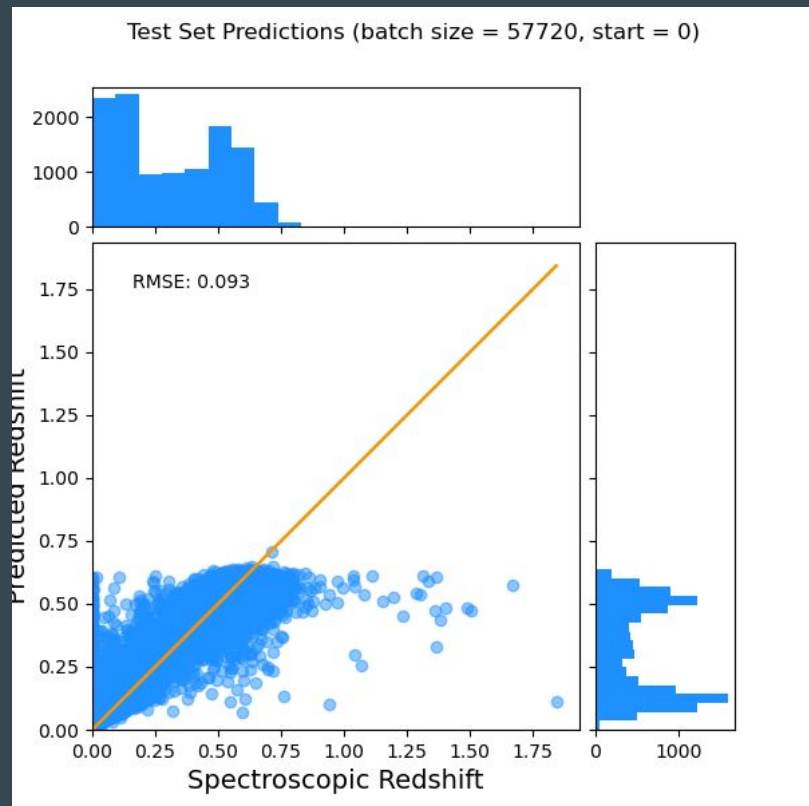
```
cv=5, scoring='neg_mean_squared_error'
```



Random Forest: Results

```
bootstrap: True  
max depth: None  
max features: None  
min samples split: 2  
n estimators: 194
```

```
'bootstrap': False,  
'max depth': None,  
'max features': sqrt,  
'min samples split': 2,  
'n_estimators': 190
```



Conclusion

Comparison of Results

Methods	RMSE
Boosted Decision Trees	0.0927
Convolutional Neural Network	0.0836
Random Forests	0.095

D'isanto and Polsterer, 2019 gives RMSE ~ 0.04

Moving Forward

- Coding best practices
 - test-driven development
 - shared repository
 - version control
 - class-based methods
 - memory optimization
- Replicability of code for robustness analysis
 - under parameter changes
 - under change in idealizing assumptions



References

- Astrophysical
 - Hoyle, *Measuring photometric redshifts using galaxy images and Deep Neural Networks*. arXiv, 2015.
 - D'Isanto and Polsterer. *Photometric redshift estimation via deep learning*. A&A 609, A111, 2018.
 - Pasquet, Bertin, et al. *Photometric redshifts from SDSS images using a Convolutional Neural Network*. A&A 621, A26, 2019.
- Machine Learning
 - Documentation and tutorials for Keras, Tensorflow, etc.
 - Towards Data Science, multiple blogs, etc.